

CWGP-AiMe vs. per-seat SaaS (e.g. toppfive.net)

Why building **AiMe** in-house beats renting a per-seat AI SaaS like toppfive.net (\$250 / \$500 / \$1000 **per user, per month**).

1. Pricing model — the big one

toppfive charges **per user, per month**. AiMe is an **internal application**: you deploy it once and **every employee uses it — there is no per-seat fee**. You pay only for what Claude actually processes (Anthropic API tokens) plus ordinary hosting.

What that means at scale (toppfive's own list prices):

Team size	toppfive Pro (\$250)	toppfive Business (\$500)	toppfive Enterprise (\$1000)	AiMe (internal)
5 users	\$1,250 / mo	\$2,500 / mo	\$5,000 / mo	API usage only
20 users	\$5,000 / mo	\$10,000 / mo	\$20,000 / mo	API usage only
50 users	\$12,500 / mo	\$25,000 / mo	\$50,000 / mo	API usage only

AiMe's cost **doesn't grow with headcount** — adding the 6th or 60th user costs nothing extra. Their bill is a **recurring rental that rises with every hire**; ours is **pay-for-actual-usage**, and the app is an asset you own. AiMe already enforces a **per-user token budget** (default 1M/period, configurable) and **rate limits**, so usage — and spend — stays predictable and controllable.

2. You own it — and can change anything

toppfive is a closed product: you get the features they ship, on their timeline. AiMe is **your codebase**. If the business needs a new capability — a bespoke RMA workflow, a NetSuite sync, a custom report — **we just build it**. No waiting for a vendor roadmap, no feature request queue, no "Enterprise-tier only" gate.

- Custom features and UI on demand.
- Custom integrations (see below).
- No vendor lock-in; the data and code are yours.

3. Native tool use via MCP — connect the tools you actually run

AiMe supports the **Model Context Protocol (MCP)**: connect a server (Slack, GitHub, Notion, ...) and the assistant **calls that tool's functions natively**, streaming, with a live "Using <tool>..." indicator. For tools without MCP, the **n8n** connector bridges to virtually anything via webhooks/workflows. That's an **open, extensible** integration story rather than a fixed connector list.

4. Your data stays internal

AiMe runs on **your** infrastructure. The Anthropic API key lives server-side and **never reaches the browser**; conversations are stored in **your** database, scoped per user. Integration secrets (n8n, MCP tokens) are **encrypted at rest**, and outbound webhook/MCP URLs are **SSRF-guarded**. With a SaaS, your prompts and business data flow through **their** platform under **their** policies.

5. Feature comparison

toppfive's advertised features, and where AiMe stands today:

toppfive feature	AiMe status
AI workspace for any business question	☐ Chat + Projects (scoped workspaces)
Document processing (PDF / CSV / Excel)	⚙ Images + PDF today; DOCX/CSV/Excel text-extraction on the roadmap
Data syncing across connected integrations	☐ via MCP (native tools) + n8n
Email monitoring & inbox triage	☐☐ Roadmap (buildable via n8n/MCP)
Scheduled workflows & recurring analysis	☐ via n8n (webhooks/schedules)
Business intelligence, reports & charts	☐☐ Roadmap
Higher query capacity / priority	☐ You control limits & model choice (no artificial tiers)
Morning briefings & daily summaries	☐☐ Roadmap (buildable via n8n schedules)
SMS alerts for critical events	☐☐ Roadmap (buildable via n8n/MCP)
Team member access	☐ Multi-user, per-user scoping
Custom workflow automation	☐ n8n + MCP + direct code changes
Dedicated support & onboarding	☐ It's your team's own app
SLA guarantees	☐ Set by your own ops, not a vendor contract

Honest take: toppfive ships some polished features AiMe hasn't built **yet** (BI dashboards, inbox triage, SMS, Office-doc extraction). The difference is that with AiMe those are a **backlog we control**, not a paywall — and several already have a clear path through n8n/MCP. Meanwhile AiMe already delivers the core (streaming chat, projects, uploads, native tools, automation, usage governance) with **no per-seat cost and full ownership**.

Bottom line

- **Cost:** their bill scales with every hire; ours scales with actual usage and never charges per seat.
- **Control:** we own the code and can add any feature or integration.
- **Privacy:** data and keys stay on our infrastructure.

- **Capability:** streaming AI chat with native MCP tools and n8n automation, today — with a roadmap we own for the rest.